



Build a 2026 AI integration playbook: map 10 agentic workflows across HR, engineering, and sales with data security and ROI scorecards

Computer

Audit your tech stack for AI integration gaps: build a vendor evaluation scorecard using NIST AI RMF and 1EdTech transparency rubrics

Computer

Show me how to map these into workflows

How do we handle institutional data and privacy

What does the student onboarding experience look like

I created a complete AI integration playbook for Semester, including 10 constrained agentic workflows, institutional data safeguards, an AI vendor scorecard, ROI measurement, and a student onboarding journey.

What the playbook covers

- 10 agentic workflows across student planning, Study Studio, academic-integrity controls, advisor preparation, campus support, career evidence, content governance, engineering reliability, RFP/sales operations, and internal talent operations.
- A reusable workflow structure defining allowed data, tools, sources, risk tier, approval gates, audit events, ROI metrics, and kill switches.
- A data-classification model for public, internal, student-private, and restricted information.
- Institutional safeguards: minimum-necessary data, tenant isolation, consent, source permissions, retention, audit logs, provider review, and no unauthorized model training on student data.
- An AI vendor evaluation scorecard aligned to NIST's Govern-Map-Measure-Manage approach and 1EdTech transparency expectations.
- A five-minute onboarding journey designed to deliver the first core outcome: "I now know what I need to do next."
- A human-confirmation matrix that prevents AI from sending, sharing, scheduling, enrolling, paying, exporting, deleting, or modifying records without an exact preview and explicit approval.
- AI quality, safety, trust, accessibility, cost, and latency metrics.

OWASP identifies prompt injection and excessive agency among major LLM application risks; the playbook therefore treats models as untrusted reasoning components, requires narrow tool permissions, validates outputs, and puts a human approval gate before consequential actions. [1842][1843]

Highest-value starting workflows

Build these first:

1. **Student Path and Registration Copilot** — explains requirements, compares courses, identifies conflicts, creates backups, and drafts advisor questions without registering students.
2. **Syllabus-to-Study Studio Agent** — turns authorized materials into source-linked study assets, reviews, practice, and study plans.
3. **Academic Integrity Guardrail Agent** — provides safe alternatives when help with active graded work is restricted.
4. **Advisor Meeting Preparation Agent** — creates a student-reviewed agenda and shares only the scope the student selects.
5. **Campus Support Navigator** — routes students to verified resources and safe official handoffs.

The playbook is available in the file panel as “semester-2026-ai-integration-playbook.”

✱✱

1. <https://genai.owasp.org/llmrisk/llm01-prompt-injection/>
2. <https://genai.owasp.org/llm-top-10/>