Semester: 16-domain deployment report

October 11, 2026 | Historical baseline | Production disconnected

Executive decision

Status: HOLD pending verified evidence. Deliverable completion is not product or launch completion. Baseline is the user screenshot and historical repository assessments, not live runtime assurance. All sixteen designed statuses are Prototype; institution-authorized and production-enabled are No in that baseline.

Source and authority boundary

Existing project authority gates require sixteen artifact categories and independent maturity facts. Prior retrieved repository flags enumerate 28 entries versus screenshot 21; count and runtime state require reconciliation. Actual hosting, migrations, queues, providers and acceptance records are unresolved. No current connector changes or live checks performed.

Six independent green gates

Designed: accepted journeys/states/roles. Implemented: authorized durable runtime. Tested: executed candidate-bound cases and qualified review. Deployed: actual target/config/schema match. Institution-authorized: exact applicable authority or accepted N/A. Production-enabled: bounded activation with staffing, recovery and reconciliation.

Owner assignments

All role assignments below are proposed. Named accountable, implementation, acceptance and backup people remain UNASSIGNED. Founder coordination does not substitute for qualified professional or institution authority.

Approval materials

Every domain requires: Product model; User screens; Data model; Workflows; Policy controls; Roles/capabilities; Audit; Support process; Monitoring/SLO; Documentation; Test plan; Release gate; Migration/rollback; Commercial model; Customer-success model; Operating owner and backup. Sensitive evidence stays in controlled storage; dashboard stores identifiers only.

Deployment runbook

Pin candidate and target manifest; rehearse additive migrations and N-1 compatibility; deploy dark with flags off; run actual-target smoke/negative checks; obtain current approvals; activate named cohort; observe accepted service and dual-run criteria. Manifest binds SHA, deployment, configuration, schema, functions/workers, objects, policy/flags, tenant/cohort/operation, approvers and expiry.

Rollback order

Declare scope -> server kill and verify direct API/worker denial -> pause/checkpoint queues -> reconcile uncertain source outcomes -> scoped rollout/cache hold -> compatible app/functions/workers rollback -> separate database/object repair -> current-revocation and post-backup reconciliation -> retest -> fresh acceptance -> bounded resume. Never wipe drafts, delete receipts, replay irreversible writes blindly, expose identities, or switch Core to a stale source.

Automation boundary

HTML evaluates local evidence failure/withdrawal/expiry and records simulated holds after evidence is entered or a local checklist was complete. Its UI refreshes every five seconds while open, not a delivery guarantee. It has no production connection. Manual trigger input represents self-reported observations. Trusted telemetry, server

enforcement, signed approvals and atomic execution still require integration.

Shared blockers

Unverified candidate/target; incomplete entry-point authorization; local/blob/duplicate authority; unwired adapters/outbox; AI bypass/broker/evaluation gaps; qualified accessibility; staffing/paging; recovery/privacy operation; institution/commercial authority; stale registers/claim conflicts.

Local verification

21 unittest cases executed and passed against the reference checker. No browser or app tests executed. The script checks consistency, not evidence authenticity or reviewer authority. Files and archive verified nonempty; PDF/PPTX export generated. No professional/customer acceptance created.

Domain scorecard and owner roles

01. Identity and onboarding

Implemented: Partial. Tested: Account-sync check. Deployed: Hosted PWA sync. Proposed roles: Identity / Security / Privacy. Named owners: UNASSIGNED.

02. Planning and degree

Implemented: Partial. Tested: Not evidenced. Deployed: Seed-data PWA. Proposed roles: Academic engineering / Registrar. Named owners: UNASSIGNED.

03. Registration

Implemented: Server; not release-ready. Tested: Server tests. Deployed: No tenant. Proposed roles: Registration / Registrar. Named owners: UNASSIGNED.

04. Courses and teaching

Implemented: Partial. Tested: Not evidenced. Deployed: Personal-course PWA. Proposed roles: Learning / Faculty. Named owners: UNASSIGNED.

05. Assignments and work

Implemented: Partial; task rows. Tested: Calendar keyboard checks. Deployed: Device-first PWA. Proposed roles: Productivity / Faculty. Named owners: UNASSIGNED.

06. Study and tutoring

Implemented: Mock + partial. Tested: Not evidenced. Deployed: Hosted PWA. Proposed roles: Learning / Tutoring operations. Named owners: UNASSIGNED.

07. Grades and progress

Implemented: Server; not release-ready. Tested: Server tests. Deployed: No tenant. Proposed roles: Gradebook / Faculty / Registrar. Named owners: UNASSIGNED.

08. Advising and support

Implemented: Partial. Tested: Not evidenced. Deployed: No staffed route. Proposed roles: Advising / Support / Privacy. Named owners: UNASSIGNED.

09. Career and portfolio

Implemented: Local/mock. Tested: Not evidenced. Deployed: Local PWA. Proposed roles: Career / Career services. Named owners: UNASSIGNED.

10. Campus and community

Implemented: Partial. Tested: Community.test.tsx. Deployed: Hosted PWA. Proposed roles: Community / Trust and Safety. Named owners: UNASSIGNED.

11. Housing, dining, services

Implemented: Partial. Tested: Not evidenced. Deployed: Hosted PWA. Proposed roles: Campus services / Providers. Named owners: UNASSIGNED.

12. Money and finance

Implemented: Partial. Tested: Billing acceptance. Deployed: Checkout held. Proposed roles: Finance / Counsel. Named owners: UNASSIGNED.

13. Institution OS

Implemented: Server; not release-ready. Tested: Server tests. Deployed: No tenant; gateway dependency. Proposed roles: Institution engineering / Customer IT. Named owners: UNASSIGNED.

14. Company OS

Implemented: Documents. Tested: Not evidenced. Deployed: Not evidenced. Proposed roles: Founder / Finance / Counsel / Operations. Named owners: UNASSIGNED.

15. Developer and partner

Implemented: None found. Tested: Not evidenced. Deployed: Not evidenced. Proposed roles: Platform / Partner operations. Named owners: UNASSIGNED.

16. Governed AI

Implemented: Partial; ungoverned. Tested: Not gated. Deployed: Shared-key function. Proposed roles: AI / Security / Privacy. Named owners: UNASSIGNED.

Twenty shared green rules

1. G-DES-01 | Designed

Journey/role/state coverage — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

2. G-DES-02 | Designed

Accessibility and consequence specification — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

3. G-IMP-01 | Implemented

UI-to-authorized-durable-state trace — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

4. G-IMP-02 | Implemented

Forged identity/tenant/owner/scope denial — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

5. G-IMP-03 | Implemented

Cross-device durability and provenance — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

6. G-TST-01 | Tested

Executed candidate-bound blocking suite — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

7. G-TST-02 | Tested

Qualified keyboard/AT/reflow acceptance — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

8. G-TST-03 | Tested

Independent security blocker closure — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

9. G-DEP-01 | Deployed

Actual build/config/schema match — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

10. G-DEP-02 | Deployed

Allowed/denied target journeys and queues — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

11. G-DEP-03 | Deployed

Previous-version rollback compatibility — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

12. G-AUT-01 | Institution-authorized

Current target/cohort/operation authority — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

13. G-AUT-02 | Institution-authorized

Withdrawn/expired approval denial — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

14. G-AUT-03 | Institution-authorized

Customer UAT and accepted limitations — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

15. G-PRD-01 | Production-enabled

Out-of-scope server/API/worker denial — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

16. G-PRD-02 | Production-enabled

Measured kill propagation — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

17. G-PRD-03 | Production-enabled

Isolated DB/object/config restore integrity — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

18. G-PRD-04 | Production-enabled

Primary/backup alert and support drill — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

19. G-PRD-05 | Production-enabled

Export/delete/revoke/offboard with holds — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

20. G-PRD-06 | Production-enabled

Operated window, outcomes and stop criteria — retain executed/accepted evidence, target binding, reviewer and validity; NOT RUN in baseline.

Eight rollback trigger classes

T01 Tenant leak or restricted disclosure

Trigger: Confirmed or credible unauthorized retrieval, including titles/search/AI context. Contain: Hold affected scope; preserve safe evidence; security/privacy incident escalation. Resume: Negative suite, impact review, remediation and fresh acceptance.

T02 Unauthorized official or financial write

Trigger: Mutation outside actor, source-system or institution authority. Contain: Kill new writeback; checkpoint commands; reconcile committed/uncertain outcomes. Resume: Source reconciliation and domain-authorized correction; fresh write authority.

T03 Corruption or reconciliation failure

Trigger: Integrity failure or unexplained divergence from authoritative source. Contain: Hold writes; isolate repair/restore; preserve audit and queues. Resume: DB/object checksums, external-state reconciliation, post-backup change review.

T04 Withdrawn or expired authority

Trigger: Applicable approval withdrawn, expired or scope-mismatched. Contain: Hold only affected operation/cohort; prevent fallback expansion. Resume: New current exact-scope approval; gate-negative retest.

T05 AI policy or gateway bypass

Trigger: Direct/proxy/worker route evades policy, budget, source access or kill. Contain: Hold relevant AI dispatch; preserve permitted audit; settle reservations. Resume: All-entrypoint enforcement, injection/kill/budget retest.

T06 Identity disclosure

Trigger: Pseudonym/private identity exposed outside accepted disclosure scope. Contain: Hold affected surface; privacy-safe transition; no name-reveal fallback. Resume: Privacy impact acceptance, corrected visibility and regression evidence.

T07 Missing support, alerts or backup

Trigger: Required coverage unavailable or alert delivery/acknowledgement fails. Contain: Hold activation/expansion; preserve urgent authorized handoff. Resume: Named coverage and successful operated alert/support drill.

T08 Accepted service threshold breached

Trigger: Populated, approved service/error/freshness/capacity threshold exceeded. Contain: Hold affected feature; diagnose against exact agreed metric/window. Resume: Recovered accepted operating window and fresh release review.